

```

In [2]: #Prac-01
#Linear regression by using Deep Neural network: Implement Boston housin
#price prediction problem by Linear regression using Deep Neural network
#Boston House price prediction dataset.
import pandas as pd
import numpy as np

# Load the Boston housing dataset from the original source
data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]

# Convert data and target into DataFrame
data_columns = ["CRIM", "ZN", "INDUS", "CHAS", "NOX",
               "RM", "AGE", "DIS", "RAD", "TAX", "PTRATIO", "B", "LSTAT"]
data = pd.DataFrame(data, columns=data_columns)
data['PRICE'] = target

# Display the first few rows of the DataFrame
print(data.head())

# Now, you can continue with the rest of the code to build and train the

```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0

	PTRATIO	B	LSTAT	PRICE
0	15.3	396.90	4.98	24.0
1	17.8	396.90	9.14	21.6
2	17.8	392.83	4.03	34.7
3	18.7	394.63	2.94	33.4
4	18.7	396.90	5.33	36.2

```

In [3]: #Check the shape of dataframe
data.shape

```

```

Out[3]: (506, 14)

```

```

In [4]: data.columns

```

```

Out[4]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD',
               'TAX',
               'PTRATIO', 'B', 'LSTAT', 'PRICE'],
              dtype='object')

```

```
In [5]: data.dtypes
```

```
Out[5]: CRIM      float64
        ZN        float64
        INDUS     float64
        CHAS      float64
        NOX       float64
        RM        float64
        AGE       float64
        DIS       float64
        RAD       float64
        TAX       float64
        PTRATIO   float64
        B        float64
        LSTAT     float64
        PRICE     float64
dtype: object
```

```
In [6]: # Identifying the unique number of values in the dataset
        data.nunique()
```

```
Out[6]: CRIM      504
        ZN        26
        INDUS     76
        CHAS      2
        NOX       81
        RM       446
        AGE      356
        DIS      412
        RAD       9
        TAX       66
        PTRATIO   46
        B       357
        LSTAT    455
        PRICE    229
dtype: int64
```

```
In [7]: # Check for missing values
        data.isnull().sum()
```

```
Out[7]: CRIM      0
        ZN        0
        INDUS     0
        CHAS      0
        NOX       0
        RM        0
        AGE       0
        DIS       0
        RAD       0
        TAX       0
        PTRATIO   0
        B        0
        LSTAT     0
        PRICE     0
dtype: int64
```

```
In [8]: # See rows with missing values
data[data.isnull().any(axis=1)]
```

```
Out[8]:
```

CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	PRICE
------	----	-------	------	-----	----	-----	-----	-----	-----	---------	---	-------	-------

```
In [9]: # Viewing the data statistics
data.describe()
```

```
Out[9]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AG
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.57490
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.14886
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.90000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.02500
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.50000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.07500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.00000



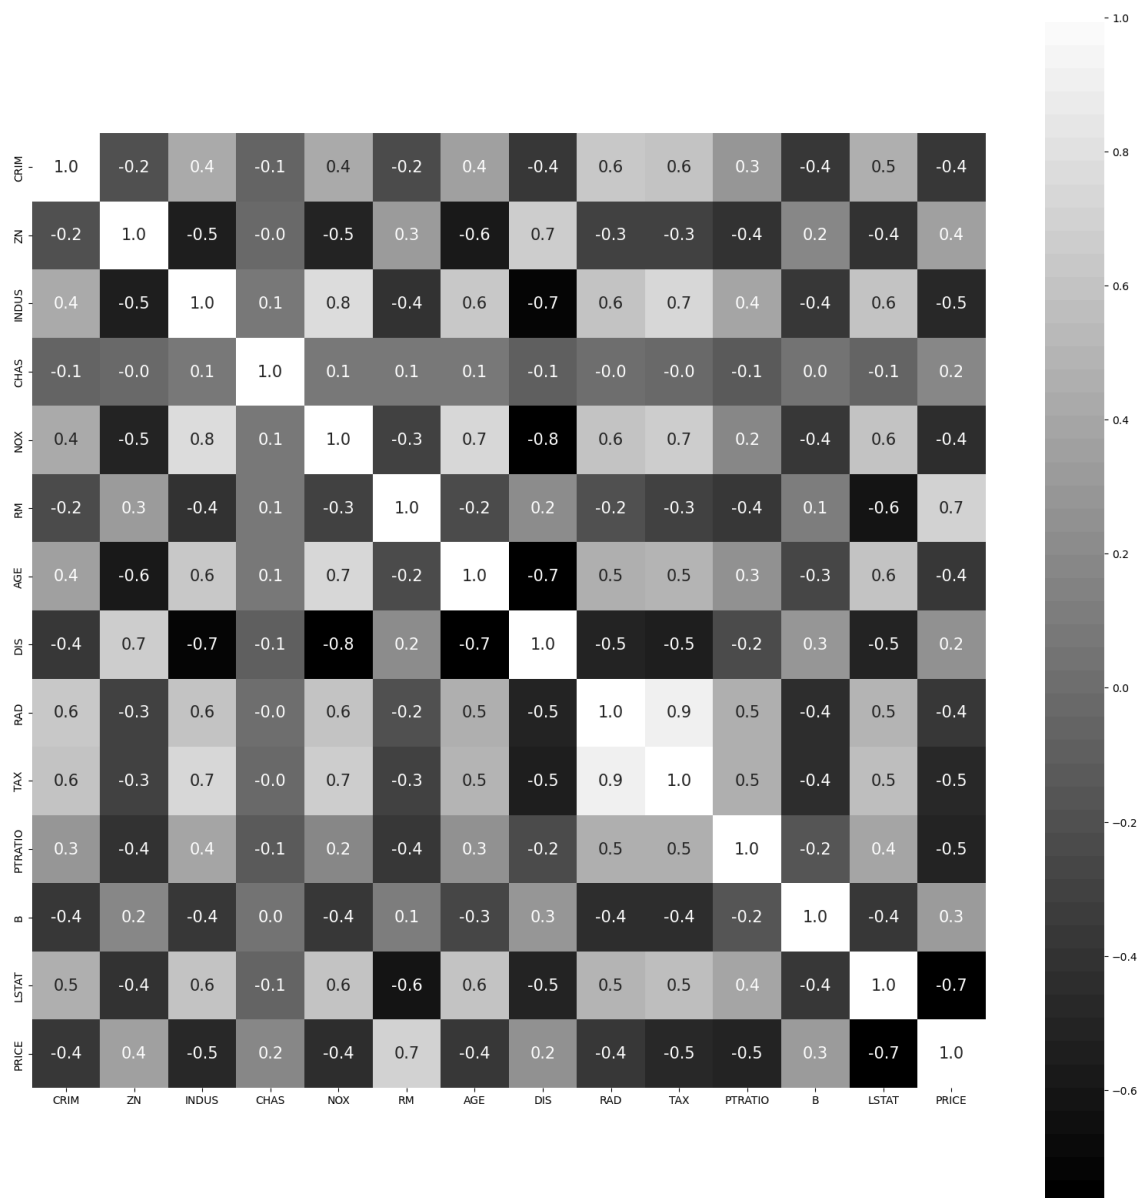
```
In [10]: # Finding out the correlation between the features
corr = data.corr()
corr.shape
```

```
Out[10]: (14, 14)
```

```
In [13]: # Plotting the heatmap of correlation between features
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(20,20))
sns.heatmap(corr, cbar=True, square=True, fmt='.1f', annot=True,
            annot_kws={'size':15}, cmap='gray')
```

Out[13]: <Axes: >



```
In [14]: # Splitting target variable and independent variables
X = data.drop(['PRICE'], axis = 1)
y = data['PRICE']
```

```
In [15]: # Splitting to training and testing data
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size = 0.3
                                                    random_state = 4)
```

```
In [16]: # Import Library for Linear Regression
from sklearn.linear_model import LinearRegression
```

```
In [17]: # Create a Linear regressor
lm = LinearRegression()
```

```
In [18]: # Train the model using the training sets
lm.fit(X_train, y_train)
```

```
Out[18]: ▾ LinearRegression
LinearRegression()
```

```
In [19]: # Value of y intercept
lm.intercept_
```

```
Out[19]: 36.35704137659398
```

```
In [21]: #Converting the coefficient values to a dataframe
coefficients = pd.DataFrame([X_train.columns, lm.coef_]).T
coefficients = coefficients.rename(columns={0: 'Attribute', 1: 'Coefficien
coefficients
```

```
Out[21]:
```

	Attribute	Coefficients
0	CRIM	-0.12257
1	ZN	0.055678
2	INDUS	-0.008834
3	CHAS	4.693448
4	NOX	-14.435783
5	RM	3.28008
6	AGE	-0.003448
7	DIS	-1.552144
8	RAD	0.32625
9	TAX	-0.014067
10	PTRATIO	-0.803275
11	B	0.009354
12	LSTAT	-0.523478

```
In [23]: from sklearn import metrics


# Model prediction on train data
y_pred = lm.predict(X_train)
# Model Evaluation
print('R^2:', metrics.r2_score(y_train, y_pred))
print('Adjusted R^2:', 1 - (1 - metrics.r2_score(y_train,
y_pred)) * (len(y_train) - 1) / (len(y_train) - X_train.shape[1] - 1))
print('MAE:', metrics.mean_absolute_error(y_train, y_pred))
print('MSE:', metrics.mean_squared_error(y_train, y_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_train, y_pred)))
```

```
R^2: 0.7465991966746854
Adjusted R^2: 0.736910342429894
MAE: 3.0898610949711265
MSE: 19.07368870346903
RMSE: 4.367343437774162
```

```
In [24]: # Visualizing the differences between actual prices and predicted values
plt.scatter(y_train, y_pred)
plt.xlabel("Prices")
plt.ylabel("Predicted prices")
plt.title("Prices vs Predicted prices")
plt.show()
```



```
In [25]: # Checking residuals
plt.scatter(y_pred,y_train-y_pred)
plt.title("Predicted vs residuals")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.show()
```



```
In [26]: # Checking Normality of errors
sns.distplot(y_train-y_pred)
plt.title("Histogram of Residuals")
plt.xlabel("Residuals")
plt.ylabel("Frequency")
plt.show()
```


C:\Users\Ekata\AppData\Local\Temp\ipykernel\_9316\3326403628.py:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

```
sns.distplot(y_train-y_pred)
```





```
In [27]: # Predicting Test data with the model
y_test_pred = lm.predict(X_test)
# Model Evaluation
acc_linreg = metrics.r2_score(y_test, y_test_pred)
print('R^2:', acc_linreg)
print('Adjusted R^2:', 1 - (1-metrics.r2_score(y_test,
y_test_pred))*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1))
print('MAE:', metrics.mean_absolute_error(y_test, y_test_pred))
print('MSE:', metrics.mean_squared_error(y_test, y_test_pred))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

```
R^2: 0.7121818377409193
Adjusted R^2: 0.6850685326005711
MAE: 3.859005592370745
MSE: 30.053993307124163
RMSE: 5.482152251362978
```

In [ ]: